Figure 3.1 Ranking skill by evaluation year

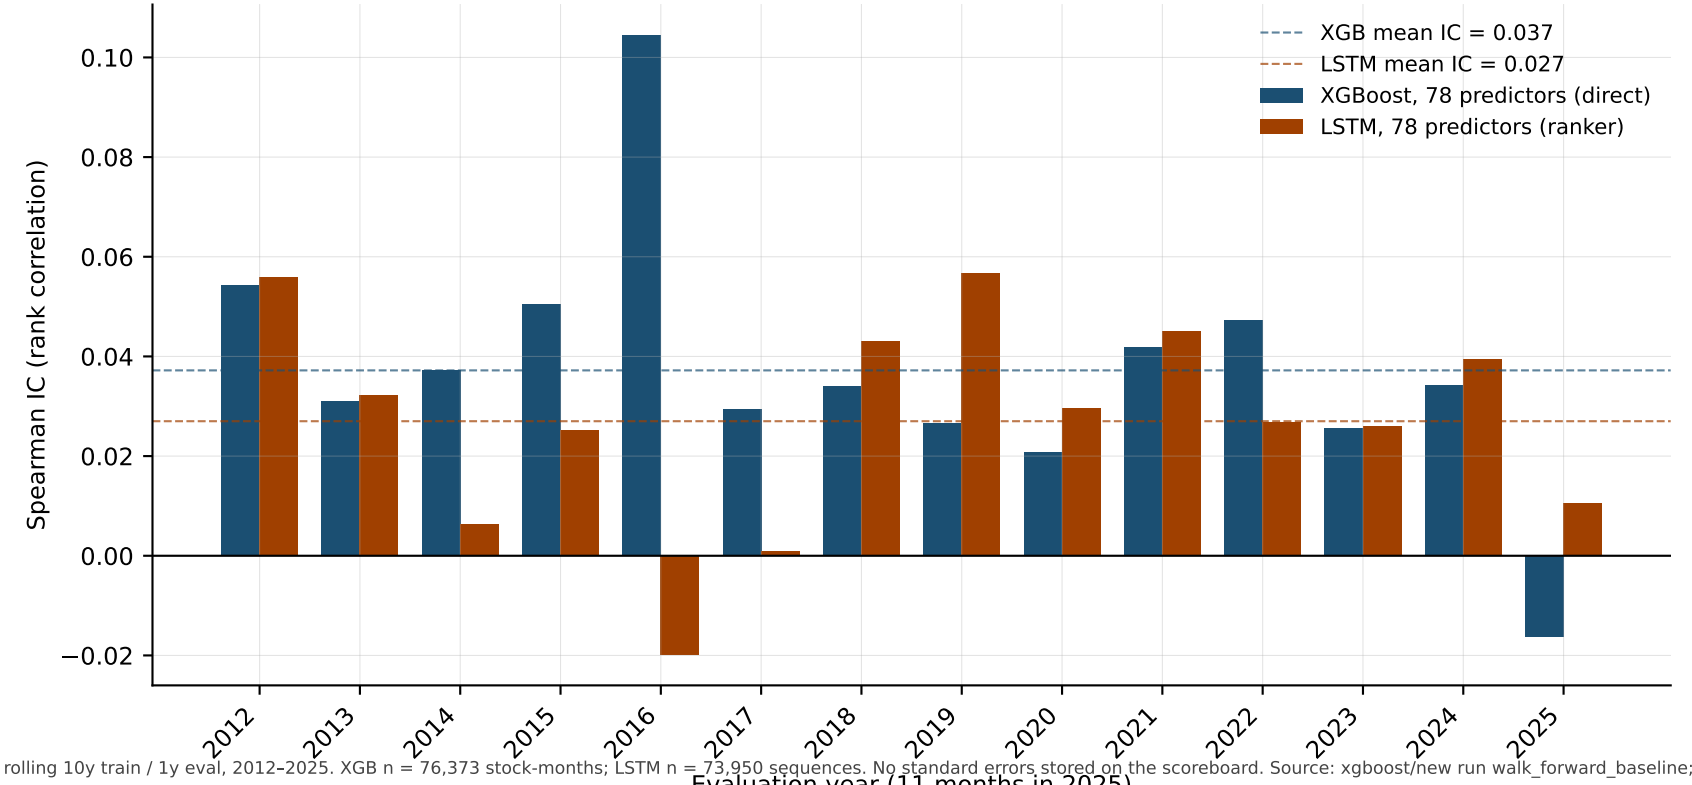

Sample: rolling 10y train / 1y eval, 2012–2025. XGB n = 76,373 stock-months; LSTM n = 73,950 sequences. No standard errors stored on the scoreboard. Source: xgboost/new run walk_forward_baseline; lstm/Full 78, 14 step No Sentiment, no golden rule.